

Hugging Face Transformers is a popular library in the field of Natural Language Processing (NLP) that provides pre-trained models and tools for various NLP tasks, including question answering (QA). The library is built on top of PyTorch and TensorFlow, and it offers an extensive collection of transformer-based models.

```
pip --version
```

```
pip 23.1.2 from /usr/local/lib/python3.10/dist-packages/pip (python 3.10)
```

```
pip install datasets
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public
Requirement already satisfied: datasets in /usr/local/lib/python3.10/dist-packages (2.12.0)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from datasets)
Requirement already satisfied: pyarrow>=8.0.0 in /usr/local/lib/python3.10/dist-packages (from datasets)
Requirement already satisfied: dill<0.3.7,>=0.3.0 in /usr/local/lib/python3.10/dist-packages (from datasets)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from datasets)
Requirement already satisfied: requests>=2.19.0 in /usr/local/lib/python3.10/dist-packages (from datasets)
Requirement already satisfied: tqdm>=4.62.1 in /usr/local/lib/python3.10/dist-packages (from datasets)
Requirement already satisfied: xxhash in /usr/local/lib/python3.10/dist-packages (from datasets)
Requirement already satisfied: multiprocessing in /usr/local/lib/python3.10/dist-packages (from datasets)
Requirement already satisfied: fsspec[http]>=2021.11.1 in /usr/local/lib/python3.10/dist-packages (from datasets)
Requirement already satisfied: aiohttp in /usr/local/lib/python3.10/dist-packages (from datasets)
Requirement already satisfied: huggingface-hub<1.0.0,>=0.11.0 in /usr/local/lib/python3.10/dist-packages (from datasets)
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from datasets)
Requirement already satisfied: responses<0.19 in /usr/local/lib/python3.10/dist-packages (from datasets)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (from datasets)
Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.10/dist-packages (from datasets)
Requirement already satisfied: charset-normalizer<4.0,>=2.0 in /usr/local/lib/python3.10/dist-packages (from datasets)
Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3.10/dist-packages (from datasets)
Requirement already satisfied: async-timeout<5.0,>=4.0.0a3 in /usr/local/lib/python3.10/dist-packages (from datasets)
Requirement already satisfied: yarl<2.0,>=1.0 in /usr/local/lib/python3.10/dist-packages (from datasets)
Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from datasets)
Requirement already satisfied: aiosignal>=1.1.2 in /usr/local/lib/python3.10/dist-packages (from datasets)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from datasets)
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.10/dist-packages (from datasets)
Requirement already satisfied: urllib3<1.27,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from datasets)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from datasets)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from datasets)
Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.10/dist-packages (from datasets)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from datasets)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from datasets)
```

```
from datasets import load_dataset
```

```
raw_datasets = load_dataset("squad")
```

```
WARNING:datasets.builder:Found cached dataset squad (/root/.cache/huggingface/datasets/squad)
100% 2/2 [00:00<00:00, 2.26it/s]
```

```
raw_datasets
```

```
DatasetDict({
  train: Dataset({
    features: ['id', 'title', 'context', 'question', 'answers'],
    num_rows: 87599
  })
  validation: Dataset({
    features: ['id', 'title', 'context', 'question', 'answers'],
    num_rows: 10570
  })
})
```

```
print("Context: ", raw_datasets["train"][0]["context"])
print("Question: ", raw_datasets["train"][0]["question"])
print("Answer: ", raw_datasets["train"][0]["answers"])
```

Context: Architecturally, the school has a Catholic character. Atop the Main Building's
Question: To whom did the Virgin Mary allegedly appear in 1858 in Lourdes France?
Answer: {'text': ['Saint Bernadette Soubirous'], 'answer_start': [515]}

```
raw_datasets["train"].filter(lambda x: len(x["answers"]["text"]) != 1)
```

WARNING:datasets.arrow_dataset:Loading cached processed dataset at /root/.cache/huggingface
Dataset({
 features: ['id', 'title', 'context', 'question', 'answers'],
 num_rows: 0
})

```
print(raw_datasets["validation"][0]["answers"])
print(raw_datasets["validation"][2]["answers"])
```

{'text': ['Denver Broncos', 'Denver Broncos', 'Denver Broncos'], 'answer_start': [177, 177, 177]}
{'text': ['Santa Clara, California', 'Levi's Stadium', 'Levi's Stadium in the San Francisco Bay Area']}

```
print(raw_datasets["validation"][2]["context"])
print(raw_datasets["validation"][2]["question"])
```

Super Bowl 50 was an American football game to determine the champion of the National Football League.
Where did Super Bowl 50 take place?

Processing the training data

```
pip install transformers
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public>
Requirement already satisfied: transformers in /usr/local/lib/python3.10/dist-packages (4.11.3)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from transformers==4.11.3) (3.12.2)
Requirement already satisfied: huggingface-hub<1.0,>=0.14.1 in /usr/local/lib/python3.10/dist-packages (from transformers==4.11.3) (0.16.4)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from transformers==4.11.3) (1.24.3)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from transformers==4.11.3) (23.1)

Requirement already satisfied: pyyaml<5.1 in /usr/local/lib/python3.10/dist-packages (from
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from t
Requirement already satisfied: tokenizers!=0.11.3,<0.14,>=0.11.1 in /usr/local/lib/python3.10/dist-packages (from transformers==4.10.0)
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.10/dist-packages (from transformers==4.10.0)
Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-packages (from huggingface-hub<0.11.0,>=0.10.0 in /usr/local/lib/python3.10/dist-packages (from transformers==4.10.0))
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.10/dist-packages (from transformers==4.10.0)
Requirement already satisfied: urllib3<1.27,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests in /usr/local/lib/python3.10/dist-packages (from transformers==4.10.0))
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests in /usr/local/lib/python3.10/dist-packages (from transformers==4.10.0))
Requirement already satisfied: charset-normalizer~=2.0.0 in /usr/local/lib/python3.10/dist-packages (from requests in /usr/local/lib/python3.10/dist-packages (from transformers==4.10.0))
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests in /usr/local/lib/python3.10/dist-packages (from transformers==4.10.0))

```
from transformers import AutoTokenizer

model_checkpoint = "bert-base-cased"
tokenizer = AutoTokenizer.from_pretrained(model_checkpoint)
```

```
tokenizer.is_fast
```

True

```
context = raw_datasets["train"][0]["context"]
question = raw_datasets["train"][0]["question"]

inputs = tokenizer(question, context)
tokenizer.decode(inputs["input_ids"])
```

'[CLS] To whom did the Virgin Mary allegedly appear in 1858 in Lourdes France? [SEP] Architecturally, the school has a Catholic character. Atop the Main Building\'s gold dome is a golden statue of the Virgin Mary. Immediately in front of the Main Building and facing it, is a copper statue of Christ with arms upraised with the legend " Venite Ad Me Omnes ". Next to the Main Building is the Basilica of the Sacred Heart. Immediately behind the basilica is the Grotto, a Marian place of prayer and reflection. It is a replica of the grotto at Lourdes, France where the Virgin Mary reputedly appeared to Saint Bernadette Soubirous in 1858. At the end of the main drive (and in a direct line that connects through 3 statues and the Gold Dome), is a simple, modern stone statue of Mary. [SEP]'

```
inputs = tokenizer(
    question,
    context,
    max_length=100,
    truncation="only_second",
    stride=50,
    return_overflowing_tokens=True,
)

for ids in inputs["input_ids"]:
    print(tokenizer.decode(ids))
```

[CLS] To whom did the Virgin Mary allegedly appear in 1858 in Lourdes France? [SEP] Architecturally, the school has a Catholic character. Atop the Main Building\'s gold dome is a golden statue of the Virgin Mary. Immediately in front of the Main Building and facing it, is a copper statue of Christ with arms upraised with the legend " Venite Ad Me Omnes ". Next to the Main Building is the Basilica of the Sacred Heart. Immediately behind the basilica is the Grotto, a Marian place of prayer and reflection. It is a replica of the grotto at Lourdes, France where the Virgin Mary reputedly appeared to Saint Bernadette Soubirous in 1858. At the end of the main drive (and in a direct line that connects through 3 statues and the Gold Dome), is a simple, modern stone statue of Mary. [SEP]

```
inputs = tokenizer(
    question,
    context,
    max_length=100,
    truncation="only_second",
    stride=50,
    return_overflowing_tokens=True,
    return_offsets_mapping=True,
)
inputs.keys()
```

```
dict_keys(['input_ids', 'token_type_ids', 'attention_mask', 'offset_mapping',
'overflow_to_sample_mapping'])
```

```
inputs["overflow_to_sample_mapping"]
```

```
[0, 0, 0, 0]
```

```
inputs = tokenizer(
    raw_datasets["train"][2:6]["question"],
    raw_datasets["train"][2:6]["context"],
    max_length=100,
    truncation="only_second",
    stride=50,
    return_overflowing_tokens=True,
    return_offsets_mapping=True,
)
```

```
print(f"The 4 examples gave {len(inputs['input_ids'])} features.")
print(f"Here is where each comes from: {inputs['overflow_to_sample_mapping']}")
```

The 4 examples gave 19 features.

Here is where each comes from: [0, 0, 0, 0, 1, 1, 1, 1, 2, 2, 2, 2, 3, 3, 3, 3, 3, 3, 3].

```
answers = raw_datasets["train"][2:6]["answers"]
start_positions = []
end_positions = []
```

```
for i, offset in enumerate(inputs["offset_mapping"]):
    sample_idx = inputs["overflow_to_sample_mapping"][i]
    answer = answers[sample_idx]
    start_char = answer["answer_start"][0]
    end_char = answer["answer_start"][0] + len(answer["text"][0])
    sequence_ids = inputs.sequence_ids(i)
```

Find the start and end of the context

```
idx = 0
```

```
while sequence_ids[idx] != 1:
```

```
    idx += 1
```

```
context_start = idx
```

```
while sequence_ids[idx] == 1:
```

```
    idx += 1
```

```
context_end = idx - 1
```

If the answer is not fully inside the context, label is (0, 0)

```
if offset[context_start][0] > start_char or offset[context_end][1] < end_char:
```

```
    start_positions.append(0)
```

```
    end_positions.append(0)
```

```
else:
```

```
    # Otherwise it's the start and end token positions
```

```
    idx = context_start
```

```

while idx <= context_end and offset[idx][0] <= start_char:
    idx += 1
start_positions.append(idx - 1)

idx = context_end
while idx >= context_start and offset[idx][1] >= end_char:
    idx -= 1
end_positions.append(idx + 1)

start_positions, end_positions

([83, 51, 19, 0, 0, 64, 27, 0, 34, 0, 0, 0, 67, 34, 0, 0, 0, 0, 0],
 [85, 53, 21, 0, 0, 70, 33, 0, 40, 0, 0, 0, 68, 35, 0, 0, 0, 0, 0])

```

```

idx = 0
sample_idx = inputs["overflow_to_sample_mapping"][idx]
answer = answers[sample_idx]["text"][0]

start = start_positions[idx]
end = end_positions[idx]
labeled_answer = tokenizer.decode(inputs["input_ids"][idx][start : end + 1])

print(f"Theoretical answer: {answer}, labels give: {labeled_answer}")

```

Theoretical answer: the Main Building, labels give: the Main Building

So that's a match


```

idx = 4
sample_idx = inputs["overflow_to_sample_mapping"][idx]
answer = answers[sample_idx]["text"][0]

decoded_example = tokenizer.decode(inputs["input_ids"][idx])
print(f"Theoretical answer: {answer}, decoded example: {decoded_example}")

```

Theoretical answer: a Marian place of prayer and reflection, decoded example: [CLS] What i



we don't see the answer inside the context

```

max_length = 384
stride = 128

def preprocess_training_examples(examples):
    questions = [q.strip() for q in examples["question"]]
    inputs = tokenizer(
        questions,
        examples["context"],
        max_length=max_length,
        truncation="only_second",
        stride=stride,
        return_overflowing_tokens=True,
        return_offsets_mapping=True,
        padding="max_length",
    )

    offset_mapping = inputs.pop("offset_mapping")

```

```

sample_map = inputs.pop("overflow_to_sample_mapping")
answers = examples["answers"]
start_positions = []
end_positions = []

for i, offset in enumerate(offset_mapping):
    sample_idx = sample_map[i]
    answer = answers[sample_idx]
    start_char = answer["answer_start"][0]
    end_char = answer["answer_start"][0] + len(answer["text"][0])
    sequence_ids = inputs.sequence_ids(i)

    # Find the start and end of the context
    idx = 0
    while sequence_ids[idx] != 1:
        idx += 1
    context_start = idx
    while sequence_ids[idx] == 1:
        idx += 1
    context_end = idx - 1

    # If the answer is not fully inside the context, label is (0, 0)
    if offset[context_start][0] > start_char or offset[context_end][1] < end_char:
        start_positions.append(0)
        end_positions.append(0)
    else:
        # Otherwise it's the start and end token positions
        idx = context_start
        while idx <= context_end and offset[idx][0] <= start_char:
            idx += 1
        start_positions.append(idx - 1)

        idx = context_end
        while idx >= context_start and offset[idx][1] >= end_char:
            idx -= 1
        end_positions.append(idx + 1)

inputs["start_positions"] = start_positions
inputs["end_positions"] = end_positions
return inputs

```

```

train_dataset = raw_datasets["train"].map(
    preprocess_training_examples,
    batched=True,
    remove_columns=raw_datasets["train"].column_names,)
len(raw_datasets["train"]), len(train_dataset)

```

WARNING:datasets.arrow_dataset:Loading cached processed dataset at /root/.cache/huggingface/datasets/87599, 88729)



the preprocessing added roughly 1,000 features

Processing the validation data

```

def preprocess_validation_examples(examples):
    questions = [q.strip() for q in examples["question"]]
    inputs = tokenizer(
        questions,

```

```

        examples["context"],
        max_length=max_length,
        truncation="only_second",
        stride=stride,
        return_overflowing_tokens=True,
        return_offsets_mapping=True,
        padding="max_length",
    )

    sample_map = inputs.pop("overflow_to_sample_mapping")
    example_ids = []

    for i in range(len(inputs["input_ids"])):
        sample_idx = sample_map[i]
        example_ids.append(examples["id"][sample_idx])

        sequence_ids = inputs.sequence_ids(i)
        offset = inputs["offset_mapping"][i]
        inputs["offset_mapping"][i] = [
            o if sequence_ids[k] == 1 else None for k, o in enumerate(offset)
        ]

    inputs["example_id"] = example_ids
    return inputs

```

```

validation_dataset = raw_datasets["validation"].map(
    preprocess_validation_examples,
    batched=True,
    remove_columns=raw_datasets["validation"].column_names,
)
len(raw_datasets["validation"]), len(validation_dataset)

```

```
(10570, 10822)
```

In this case we've only added a couple of hundred samples, so it appears the contexts in the validation dataset are a bit shorter.

Fine-tuning the model with the Trainer API

Post Processing


```

small_eval_set = raw_datasets["validation"].select(range(100))
trained_checkpoint = "distilbert-base-cased-distilled-squad"

tokenizer = AutoTokenizer.from_pretrained(trained_checkpoint)
eval_set = small_eval_set.map(
    preprocess_validation_examples,
    batched=True,
    remove_columns=raw_datasets["validation"].column_names,
)

```

```
WARNING:datasets.arrow_dataset:Loading cached processed dataset at /root/.cache/huggingface
```



```
tokenizer = AutoTokenizer.from_pretrained(model_checkpoint)
```

```

import torch
from transformers import AutoModelForQuestionAnswering

eval_set_for_model = eval_set.remove_columns(["example_id", "offset_mapping"])
eval_set_for_model.set_format("torch")

device = torch.device("cuda") if torch.cuda.is_available() else torch.device("cpu")
batch = {k: eval_set_for_model[k].to(device) for k in eval_set_for_model.column_names}
trained_model = AutoModelForQuestionAnswering.from_pretrained(trained_checkpoint).to(
    device)

with torch.no_grad():
    outputs = trained_model(**batch)

```

Since the Trainer will give us predictions as NumPy arrays, we grab the start and end logits and convert them to that format

```

start_logits = outputs.start_logits.cpu().numpy()
end_logits = outputs.end_logits.cpu().numpy()

```

```

import collections

example_to_features = collections.defaultdict(list)
for idx, feature in enumerate(eval_set):
    example_to_features[feature["example_id"]].append(idx)

```

```

import numpy as np

n_best = 20
max_answer_length = 30
predicted_answers = []

for example in small_eval_set:
    example_id = example["id"]
    context = example["context"]
    answers = []

    for feature_index in example_to_features[example_id]:
        start_logit = start_logits[feature_index]
        end_logit = end_logits[feature_index]
        offsets = eval_set["offset_mapping"][feature_index]

        start_indexes = np.argsort(start_logit)[-1 : -n_best - 1 : -1].tolist()
        end_indexes = np.argsort(end_logit)[-1 : -n_best - 1 : -1].tolist()
        for start_index in start_indexes:
            for end_index in end_indexes:
                # Skip answers that are not fully in the context
                if offsets[start_index] is None or offsets[end_index] is None:
                    continue
                # Skip answers with a length that is either < 0 or > max_answer_length.
                if (
                    end_index < start_index
                    or end_index - start_index + 1 > max_answer_length
                ):
                    continue

                answers.append(
                    {
                        "text": context[offsets[start_index][0] : offsets[end_index][1]],
                        "logit_score": start_logit[start_index] + end_logit[end_index],

```



```

    }
)

best_answer = max(answers, key=lambda x: x["logit_score"])
predicted_answers.append({"id": example_id, "prediction_text": best_answer["text"]})

```

```
pip install evaluate
```

```

Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public
Requirement already satisfied: evaluate in /usr/local/lib/python3.10/dist-packages (0.4.0)
Requirement already satisfied: datasets>=2.0.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from evaluate)
Requirement already satisfied: dill in /usr/local/lib/python3.10/dist-packages (from evaluate)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from evaluate)
Requirement already satisfied: requests>=2.19.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: tqdm>=4.62.1 in /usr/local/lib/python3.10/dist-packages (from evaluate)
Requirement already satisfied: xxhash in /usr/local/lib/python3.10/dist-packages (from evaluate)
Requirement already satisfied: multiprocessing in /usr/local/lib/python3.10/dist-packages (from evaluate)
Requirement already satisfied: fsspec[http]>=2021.05.0 in /usr/local/lib/python3.10/dist-packages (from evaluate)
Requirement already satisfied: huggingface-hub>=0.7.0 in /usr/local/lib/python3.10/dist-packages (from evaluate)
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from evaluate)
Requirement already satisfied: responses<0.19 in /usr/local/lib/python3.10/dist-packages (from evaluate)
Requirement already satisfied: pyarrow>=8.0.0 in /usr/local/lib/python3.10/dist-packages (from evaluate)
Requirement already satisfied: aiohttp in /usr/local/lib/python3.10/dist-packages (from evaluate)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (from evaluate)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from evaluate)
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.10/dist-packages (from evaluate)
Requirement already satisfied: urllib3<1.27,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from evaluate)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from evaluate)
Requirement already satisfied: charset-normalizer~=2.0.0 in /usr/local/lib/python3.10/dist-packages (from evaluate)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from evaluate)
Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.10/dist-packages (from evaluate)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from evaluate)
Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.10/dist-packages (from evaluate)
Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3.10/dist-packages (from evaluate)
Requirement already satisfied: async-timeout<5.0,>=4.0.0a3 in /usr/local/lib/python3.10/dist-packages (from evaluate)
Requirement already satisfied: yarll<2.0,>=1.0 in /usr/local/lib/python3.10/dist-packages (from evaluate)
Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from evaluate)
Requirement already satisfied: aiosignal>=1.1.2 in /usr/local/lib/python3.10/dist-packages (from evaluate)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from evaluate)

```

```
import evaluate
```

```
metric = evaluate.load("squad")
```

```

theoretical_answers = [
    {"id": ex["id"], "answers": ex["answers"]} for ex in small_eval_set
]

```

```

print(predicted_answers[0])
print(theoretical_answers[0])

```

```

{'id': '56be4db0acb8001400a502ec', 'prediction_text': 'Denver Broncos'}
{'id': '56be4db0acb8001400a502ec', 'answers': {'text': ['Denver Broncos', 'Denver Broncos']}}

```

```
metric.compute(predictions=predicted_answers, references=theoretical_answers)
```

```
{'exact_match': 83.0, 'f1': 88.25000000000004}
```

Now let's put everything we just did in a `compute_metrics()` function that we will use in the Trainer

```
from tqdm.auto import tqdm

def compute_metrics(start_logits, end_logits, features, examples):
    example_to_features = collections.defaultdict(list)
    for idx, feature in enumerate(features):
        example_to_features[feature["example_id"]].append(idx)

    predicted_answers = []
    for example in tqdm(examples):
        example_id = example["id"]
        context = example["context"]
        answers = []

        # Loop through all features associated with that example
        for feature_index in example_to_features[example_id]:
            start_logit = start_logits[feature_index]
            end_logit = end_logits[feature_index]
            offsets = features[feature_index]["offset_mapping"]

            start_indexes = np.argsort(start_logit)[-1 : -n_best - 1 : -1].tolist()
            end_indexes = np.argsort(end_logit)[-1 : -n_best - 1 : -1].tolist()
            for start_index in start_indexes:
                for end_index in end_indexes:
                    # Skip answers that are not fully in the context
                    if offsets[start_index] is None or offsets[end_index] is None:
                        continue
                    # Skip answers with a length that is either < 0 or > max_answer_length
                    if (
                        end_index < start_index
                        or end_index - start_index + 1 > max_answer_length
                    ):
                        continue

                    answer = {
                        "text": context[offsets[start_index][0] : offsets[end_index][1]],
                        "logit_score": start_logit[start_index] + end_logit[end_index],
                    }
                    answers.append(answer)

        # Select the answer with the best score
        if len(answers) > 0:
            best_answer = max(answers, key=lambda x: x["logit_score"])
            predicted_answers.append(
                {"id": example_id, "prediction_text": best_answer["text"]}
            )
        else:
            predicted_answers.append({"id": example_id, "prediction_text": ""})

    theoretical_answers = [{"id": ex["id"], "answers": ex["answers"]} for ex in examples]
    return metric.compute(predictions=predicted_answers, references=theoretical_answers)
```

```
compute_metrics(start_logits, end_logits, eval_set, small_eval_set)
```

100%

100/100 [00:00<00:00, 619.69it/s]

Fine-tuning the model

```
model = AutoModelForQuestionAnswering.from_pretrained(model_checkpoint)
```

Some weights of the model checkpoint at bert-base-cased were not used when initializing BertForQuestionAnswering
- This IS expected if you are initializing BertForQuestionAnswering from the checkpoint of bert-base-cased
- This IS NOT expected if you are initializing BertForQuestionAnswering from the checkpoint of bert-base-cased
Some weights of BertForQuestionAnswering were not initialized from the model checkpoint at bert-base-cased
You should probably TRAIN this model on a down-stream task to be able to use it for predictions.

```
from huggingface_hub import notebook_login
```

```
notebook_login()
```

Token is valid (permission: write).

Your token has been saved in your configured git credential helpers

Your token has been saved to /root/.cache/huggingface/token

Login successful

```
pip install --upgrade accelerate
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public>
Requirement already satisfied: accelerate in /usr/local/lib/python3.10/dist-packages (0.19.0)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from accelerate==0.19.0)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from accelerate==0.19.0)
Requirement already satisfied: psutil in /usr/local/lib/python3.10/dist-packages (from accelerate==0.19.0)
Requirement already satisfied: pyyaml in /usr/local/lib/python3.10/dist-packages (from accelerate==0.19.0)
Requirement already satisfied: torch>=1.6.0 in /usr/local/lib/python3.10/dist-packages (from accelerate==0.19.0)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from torch>=1.6.0 to accelerate==0.19.0)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.10/dist-packages (from torch>=1.6.0 to accelerate==0.19.0)
Requirement already satisfied: sympy in /usr/local/lib/python3.10/dist-packages (from torch>=1.6.0 to accelerate==0.19.0)
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (from torch>=1.6.0 to accelerate==0.19.0)
Requirement already satisfied: jinja2 in /usr/local/lib/python3.10/dist-packages (from torch>=1.6.0 to accelerate==0.19.0)
Requirement already satisfied: triton==2.0.0 in /usr/local/lib/python3.10/dist-packages (from torch>=1.6.0 to accelerate==0.19.0)
Requirement already satisfied: cmake in /usr/local/lib/python3.10/dist-packages (from triton==2.0.0 to torch>=1.6.0 to accelerate==0.19.0)
Requirement already satisfied: lit in /usr/local/lib/python3.10/dist-packages (from triton==2.0.0 to torch>=1.6.0 to accelerate==0.19.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from jinja2 to torch>=1.6.0 to accelerate==0.19.0)
Requirement already satisfied: mpmath>=0.19 in /usr/local/lib/python3.10/dist-packages (from sympy to torch>=1.6.0 to accelerate==0.19.0)

```
pip install git+https://github.com/huggingface/accelerate
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public>
Collecting git+<https://github.com/huggingface/accelerate>
Cloning <https://github.com/huggingface/accelerate> to /tmp/pip-req-build-v6op_4ny
Running command git clone --filter=blob:none --quiet <https://github.com/huggingface/accelerate>
Resolved <https://github.com/huggingface/accelerate> to commit 85901cdcf99e9fd258811789ca6
Installing build dependencies ... done
Getting requirements to build wheel ... done
Preparing metadata (pyproject.toml) ... done
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from accelerate==0.20.0.dev0)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from accelerate==0.20.0.dev0)
Requirement already satisfied: psutil in /usr/local/lib/python3.10/dist-packages (from accelerate==0.20.0.dev0)
Requirement already satisfied: pyyaml in /usr/local/lib/python3.10/dist-packages (from accelerate==0.20.0.dev0)
Requirement already satisfied: torch>=1.6.0 in /usr/local/lib/python3.10/dist-packages (from accelerate==0.20.0.dev0)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from torch>=1.6.0 to accelerate==0.20.0.dev0)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.10/dist-packages (from torch>=1.6.0 to accelerate==0.20.0.dev0)
Requirement already satisfied: sympy in /usr/local/lib/python3.10/dist-packages (from torch>=1.6.0 to accelerate==0.20.0.dev0)
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (from torch>=1.6.0 to accelerate==0.20.0.dev0)
Requirement already satisfied: Jinja2 in /usr/local/lib/python3.10/dist-packages (from torch>=1.6.0 to accelerate==0.20.0.dev0)
Requirement already satisfied: triton==2.0.0 in /usr/local/lib/python3.10/dist-packages (from torch>=1.6.0 to accelerate==0.20.0.dev0)
Requirement already satisfied: cmake in /usr/local/lib/python3.10/dist-packages (from triton==2.0.0 to torch>=1.6.0 to accelerate==0.20.0.dev0)
Requirement already satisfied: lit in /usr/local/lib/python3.10/dist-packages (from triton==2.0.0 to torch>=1.6.0 to accelerate==0.20.0.dev0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from Jinja2 to torch>=1.6.0 to accelerate==0.20.0.dev0)
Requirement already satisfied: mpmath>=0.19 in /usr/local/lib/python3.10/dist-packages (from sympy to torch>=1.6.0 to accelerate==0.20.0.dev0)
Building wheels for collected packages: accelerate
Building wheel for accelerate (pyproject.toml) ... done
Created wheel for accelerate: filename=accelerate-0.20.0.dev0-py3-none-any.whl size=226911 sha256=85901cdcf99e9fd258811789ca6
Stored in directory: /tmp/pip-ephem-wheel-cache-na6dr0au/wheels/f6/c7/9d/1b8a5ca8353d936
Successfully built accelerate
Installing collected packages: accelerate
Attempting uninstall: accelerate
Found existing installation: accelerate 0.19.0
Uninstalling accelerate-0.19.0:
Successfully uninstalled accelerate-0.19.0

```
pip uninstall -y transformers accelerate
```

Found existing installation: transformers 4.28.0
Uninstalling transformers-4.28.0:
Successfully uninstalled transformers-4.28.0
Found existing installation: accelerate 0.20.0.dev0
Uninstalling accelerate-0.20.0.dev0:
Successfully uninstalled accelerate-0.20.0.dev0

```
pip install transformers accelerate
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public>
 Collecting transformers
 Using cached transformers-4.29.2-py3-none-any.whl (7.1 MB)
 Collecting accelerate
 Using cached accelerate-0.19.0-py3-none-any.whl (219 kB)
 Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from transformers==4.29.2)
 Requirement already satisfied: huggingface-hub<1.0,>=0.14.1 in /usr/local/lib/python3.10/dist-packages (from transformers==4.29.2)
 Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from transformers==4.29.2)
 Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from transformers==4.29.2)
 Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (from transformers==4.29.2)
 Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.10/dist-packages (from transformers==4.29.2)
 Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from transformers==4.29.2)
 Requirement already satisfied: tokenizers!=0.11.3,<0.14,>=0.11.1 in /usr/local/lib/python3.10/dist-packages (from transformers==4.29.2)
 Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.10/dist-packages (from transformers==4.29.2)
 Requirement already satisfied: psutil in /usr/local/lib/python3.10/dist-packages (from accelerate==0.19.0)
 Requirement already satisfied: torch>=1.6.0 in /usr/local/lib/python3.10/dist-packages (from accelerate==0.19.0)
 Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-packages (from huggingface-hub<1.0,>=0.14.1->transformers==4.29.2)
 Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub<1.0,>=0.14.1->transformers==4.29.2)
 Requirement already satisfied: sympy in /usr/local/lib/python3.10/dist-packages (from torch>=1.6.0->accelerate==0.19.0)
 Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (from torch>=1.6.0->accelerate==0.19.0)
 Requirement already satisfied: jinja2 in /usr/local/lib/python3.10/dist-packages (from torch>=1.6.0->accelerate==0.19.0)

```
from transformers import TrainingArguments
```

```
args = TrainingArguments(
    "bert-finetuned-squad",
    evaluation_strategy="no",
    save_strategy="epoch",
    learning_rate=2e-5,
    num_train_epochs=3,
    weight_decay=0.01,
    # fp16=True,
    push_to_hub=True)
```

[accelerate, transformers]

```
from transformers import Trainer
```

```
trainer = Trainer(
    model=model,
    args=args,
    train_dataset=train_dataset,
    eval_dataset=validation_dataset,
    tokenizer=tokenizer,)
trainer.train()
```


Cloning <https://huggingface.co/srish884/bert-finetuned-squad> into local empty directory.
 WARNING:huggingface_hub.repository:Cloning <https://huggingface.co/srish884/bert-finetuned-squad> into local empty directory.
 /usr/local/lib/python3.10/dist-packages/transformers/optimization.py:407: FutureWarning: 1
 warnings.warn(
 You're using a BertTokenizerFast tokenizer. Please note that with a fast tokenizer, using

[124/33276 50:03 < 226:43:56, 0.04 it/s, Epoch 0.01/3]

Step Training Loss

[235/33276 1:36:22 < 227:46:28, 0.04 it/s, Epoch 0.02/3]

Step Training Loss



The whole Training takes over an hour

```
# predictions, _, _ = trainer.predict(validation_dataset)
# start_logits, end_logits = predictions
# compute_metrics(start_logits, end_logits, validation_dataset, raw_datasets["validation"])
```

```
# trainer.push_to_hub(commit_message="Training complete")
```

```
# from torch.utils.data import DataLoader
# from transformers import default_data_collator

# train_dataset.set_format("torch")
# validation_set = validation_dataset.remove_columns(["example_id", "offset_mapping"])
# validation_set.set_format("torch")

# train_dataloader = DataLoader(
#     train_dataset,
#     shuffle=True,
#     collate_fn=default_data_collator,
#     batch_size=8,
# )
# eval_dataloader = DataLoader(
#     validation_set, collate_fn=default_data_collator, batch_size=8
# )
```

```
# model = AutoModelForQuestionAnswering.from_pretrained(model_checkpoint)
```

```
# from torch.optim import AdamW

# optimizer = AdamW(model.parameters(), lr=2e-5)
```

```
# from accelerate import Accelerator

# accelerator = Accelerator(fp16=True)
# model, optimizer, train_dataloader, eval_dataloader = accelerator.prepare(
#     model, optimizer, train_dataloader, eval_dataloader
# )
```

```
# from transformers import get_scheduler

# num_train_epochs = 3
# num_update_steps_per_epoch = len(train_dataloader)
# num_training_steps = num_train_epochs * num_update_steps_per_epoch

# lr_scheduler = get_scheduler(
#     "linear",
#     optimizer=optimizer,
#     num_warmup_steps=0,
#     num_training_steps=num_training_steps,
# )
```

```
# from transformers import get_scheduler

# num_train_epochs = 3
# num_update_steps_per_epoch = len(train_dataloader)
# num_training_steps = num_train_epochs * num_update_steps_per_epoch

# lr_scheduler = get_scheduler(
#     "linear",
#     optimizer=optimizer,
```

```
# num_warmup_steps=0,  
# num_training_steps=num_training_steps)
```

```
# output_dir = "bert-finetuned-squad-accelerate"  
# repo = Repository(output_dir, clone_from=repo_name)
```

The complete code for the training loop

```
# from tqdm.auto import tqdm  
# import torch  
  
# progress_bar = tqdm(range(num_training_steps))  
  
# for epoch in range(num_train_epochs):  
#     # Training  
#     model.train()  
#     for step, batch in enumerate(train_dataloader):  
#         outputs = model(**batch)  
#         loss = outputs.loss  
#         accelerator.backward(loss)  
  
#         optimizer.step()  
#         lr_scheduler.step()  
#         optimizer.zero_grad()  
#         progress_bar.update(1)  
  
#     # Evaluation  
#     model.eval()  
#     start_logits = []  
#     end_logits = []  
#     accelerator.print("Evaluation!")  
#     for batch in tqdm(eval_dataloader):  
#         with torch.no_grad():  
#             outputs = model(**batch)  
  
#             start_logits.append(accelerator.gather(outputs.start_logits).cpu().numpy())  
#             end_logits.append(accelerator.gather(outputs.end_logits).cpu().numpy())  
  
#     start_logits = np.concatenate(start_logits)  
#     end_logits = np.concatenate(end_logits)  
#     start_logits = start_logits[: len(validation_dataset)]  
#     end_logits = end_logits[: len(validation_dataset)]  
  
#     metrics = compute_metrics(  
#         start_logits, end_logits, validation_dataset, raw_datasets["validation"]  
#     )  
#     print(f"epoch {epoch}:", metrics)  
  
#     # Save and upload  
#     accelerator.wait_for_everyone()  
#     unwrapped_model = accelerator.unwrap_model(model)  
#     unwrapped_model.save_pretrained(output_dir, save_function=accelerator.save)  
#     if accelerator.is_main_process:  
#         tokenizer.save_pretrained(output_dir)  
#         repo.push_to_hub(  
#             commit_message=f"Training in progress epoch {epoch}", blocking=False  
#         )  
  
# accelerator.wait_for_everyone()  
# unwrapped_model = accelerator.unwrap_model(model)
```

```
# unwrapped_model.save_pretrained(output_dir, save_function=accelerator.save)
```

```
from transformers import pipeline
```

```
model_checkpoint = "https://huggingface.co/srish884/bert-finetuned-squad"  
question_answerer = pipeline("question-answering", model=model_checkpoint)
```

```
context = ""
```

```
HuggingFace Transformers is backed by the three most popular deep learning libraries – Jax, PyTorch an  
between them. It's straightforward to train your models with one before loading them for inference wit  
""
```

```
question = "Which deep learning libraries back HuggingFace Transformers?"  
question_answerer(question=question, context=context)
```

```
{'score': 0.9979003071784973,  
  'start': 78,  
  'end': 105,  
  'answer': 'Jax, PyTorch and TensorFlow'}
```